

---

# CENG 352

## Database Management Systems

### Spring 2024-2025

### Project 1

---

April 13, 2025, Sunday, 23:59

## 1. Objectives

In this project, you are asked to write several SQL queries on a relational database using the given dataset. Your SQL server will be PostgreSQL, and you will work on your local environment. This project consists of 4 parts:

- Creating the database using the given CSV files.
- Writing proper SQL queries for certain problems.
- Testing the effect of indexing on query execution.
- Creating triggers and views.

## 2. Database Schema

### 2.1 Tables

Here are the database tables and their columns:

<b>Country</b> name code capital area population
---

<b>City</b> name country_code population elevation
<b>Economy</b> country_code gdp agriculture industry service inflation unemployment
<i>percentage of services of the GDP</i> <i>percentage of services of the GDP</i> <i>percentage of services of the GDP</i>
<b>Religion</b> country_code name percentage
<b>Spoken</b> country_code language percentage
<b>Encompasses</b> country_code continent_name percentage
<b>Continent</b> name area

### 3. Tasks

#### 3.1 Task 1 - Creating the Database (15 pts)

Using the given CSV files and considering the database schema above, you should create a file named `task1.sql` that contains SQL statements that you've used to create the tables. You need to use PostgreSQL.

## 3.2 Task 2 - SQL Queries (50 pts)

For this task, you should be able to write SQL queries for given problems. See the given query results to understand the structure. Create an individual SQL file for each subtask. Name the SQL files in this format: `task2_1.sql`, `task2_2.sql`, etc. Each task is worth 10 points.

### 3.2.1 Regulations

- For 3rd question, a country is inside a continent if encompasses table percentage of the country-continent tuple is more than all other country-continent tuples for that country.
- For 4th question, a country is inside a continent if encompasses table percentage of the country-continent tuple is more than 50%.
- For 5th question, a country can be in multiple continents so if a country of 1000 people is 40% in Europe, assume 400 people of that country contributes to the total population of Europe

### 3.2.2 Questions

1. Identify cities with population exceeding 50% of their country's average city population using PARTITION BY.
2. Divide countries into 4 equal groups based on their total area. Return quartile number, number of countries in each quartile, and the average area for that quartile. (*Hint: NTILE function*)
3. Find top 3 continents by GDP per capita. Return continent name, total GDP, avg GDP per capita, ordered by avg GDP per capita descending.
4. Write a single SQL query that computes the statistics present in a cross tabulation over continent and sectors(agriculture/industry/service). The aggregate value reported should be total gdp for each continent-sector pair. Each row should contain continent, sector, total GDP. You must solve this using CUBE function.
5. Use GROUPING SETS to show world/continent aggregates(Only use Encompasses and Country tables Don't use the Continent table):
  - Total area,
  - Total population

## 3.3 Task 3 - Index Optimization (20 pts)

Make the following queries. Run as fast as possible by adding indexes to the tables. Your goal is to find the least estimated execution time according to EXPLAIN statement in sql. To do so, you may create BTREE and HASH indexes. You may create these on any table, and with any set of columns. Write your answers in txt files: `task3_1.txt`, `task3_2.txt`, ...

For full credit, your answer to each problem must have the following four pieces of information:

1. CREATE INDEX statements showing which indexes you propose to use
2. An explanation for why you chose each of the indexes (e.g. the query is looking for a range of values, so we use BTREE index instead of a HASH index)
3. The output of EXPLAIN for the query after the indexes were added
4. Whether the indexes caused a performance increase for the query (just state if there was a speed up or not. You do not need to speculate why the index(es) did not speed up the query.)

To check the performance of each query independently, you should delete any indexes before moving on to the next problem. You can do so using DROP INDEX

## Questions

Q1

```
SELECT c.name, ct.name, c.population, e.gdp, e.inflation
FROM city c
NATURAL JOIN country ct
NATURAL JOIN economy e
```

Q2

```
SELECT c.name, ct.name, c.population, e.gdp, e.inflation
FROM city c
NATURAL JOIN country ct
NATURAL JOIN economy e
WHERE c.name LIKE 'B%'
      AND ct.name = 'Turkey'
```

Q3

```
SELECT c.name, ct.name, c.population, e.gdp, e.inflation
FROM city c
NATURAL JOIN country ct
NATURAL JOIN economy e
WHERE e.gdp > 100000
```

Q4

```
SELECT c.name, ct.name, c.population, e.gdp, e.inflation
FROM city c
NATURAL JOIN country ct
NATURAL JOIN economy e
WHERE c.population > 5000000
      AND inflation = 4
ORDER BY population
```

**IMPORTANT:** Many databases, such as MySQL and PostgreSQL, automatically create indexes on primary keys and unique constraints and foreign keys. If a table has a foreign key constraint, an index might already exist on that column, making manual index creation redundant. You can check these existing indexes via.

```
SELECT tablename, indexname, indexdef
FROM pg_indexes
WHERE schemaname = 'public'
order by tablename;
```

### 3.4 Task 4 - Views and Triggers (15 pts)

#### 3.4.1 View(5pts)

Create a view for:

- Country Name
- Biggest City(In population)
- Highest Elevation
- Density(Population/Area)

If a country does not have any cities listed in the city table, display **NULL** for both the largest city and the highest elevation.

#### 3.4.2 Aggregated Table(5pts)

Create a Table named **country\_city\_count** whose columns can be seen below. And populate it. Your sql file submission should contain the create and insert statements.

- Country Name
- Country Code
- Number of Cities(inside the country)

#### 3.4.3 Triggers(5pts)

Implement two triggers to update the table above when:

- A new city is inserted.
- A city's occupation changes (i.e., its country is updated).

## 4. Regulations

### 4.1 Submission

Submission will be done via ODTUClass. Please remember that **late submission is allowed up to 5 days for the entire semester**. You can distribute these 5 days to any project you want. You should put the answer query to each question in a separate .sql file and tar those .sql files with the following name: `e1234567_project1.tar.gz`

- `task1.sql`
- `task2_1.sql`
- `task2_2.sql`
- ...
- `task2_5.sql`
- `task3_1.txt`
- `task3_2.txt`
- `task3_3.txt`
- `task3_4.txt`
- `task4_1.sql`
- `task4_2.sql`
- `task4_3.sql`

Where you should replace "1234567" with your own student number.

### 4.2 SQL Style

You should write your queries in a readable manner. Write each clause on a separate line, add indentation for a clear look. For example:

```
SELECT *  
FROM orders o  
WHERE o.order_id = 'test' AND  
      o.testing_phase = 1
```

is easier to read than:

```
SELECT * FROM orders o WHERE o.order_id = 'test' AND o.testing_phase = 1
```

You can use online formatters/beautifiers for better indentation if you want.

## 4.3 Newsgroup

You must follow the ODTUClass for discussions and possible updates on a daily basis.

## 5. Tutorial & Guides

- You can download PostgreSQL server from **here**. Also an installation document for ubuntu users have been uploaded with this document.
- For visualization, you can use **DBeaver** or **DataGrip**. You can also use them for many other database servers. **DBeaver** guide is given below **DataGrip** guide is given in another document.
- Once you start the database server, open DBeaver and click 'New Database Connection' on the top left corner. (Figure 1)
- Choose PostgreSQL. (Figure 2)
- Enter credentials and connect to the default 'postgres' database. You will create your own database for this project later. (Figure 3)
- Open a new script and execute "create database <database\_name>" by selecting the script and hitting (CTRL + ENTER). (Figure 4)
- Now that you have created another database, connect to the newly created database just like before. (Figure 5)
- Open 'Tables' to see the tables. (Figure 6) Since there are no tables, you need to create tables using your script. Run queries by selecting parts of the script and hitting CTRL + ENTER. (Figure 7)
- You are almost there. Now, you need to import the data to tables.
- To import data, right-click on the table name and click 'Import data' (Figure 8)
- Choose CSV type and .csv file that you want to import for the selected table. (Figure 9, 10)
- Choose **comma(,)** as the delimiter for the given CSV files. You can observe that in individual CSV files.
- That's it. Now you can write queries on tables.
- Note: Don't mind that the figures are old, the import behavior is exactly the same.




Figure 1: New database connection

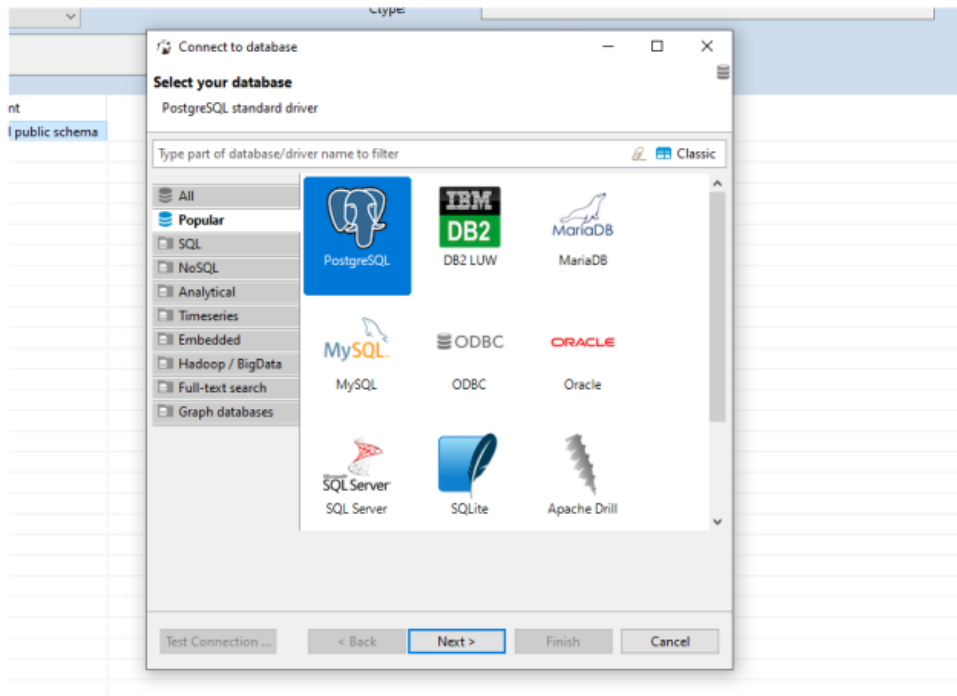


Figure 2: Database options



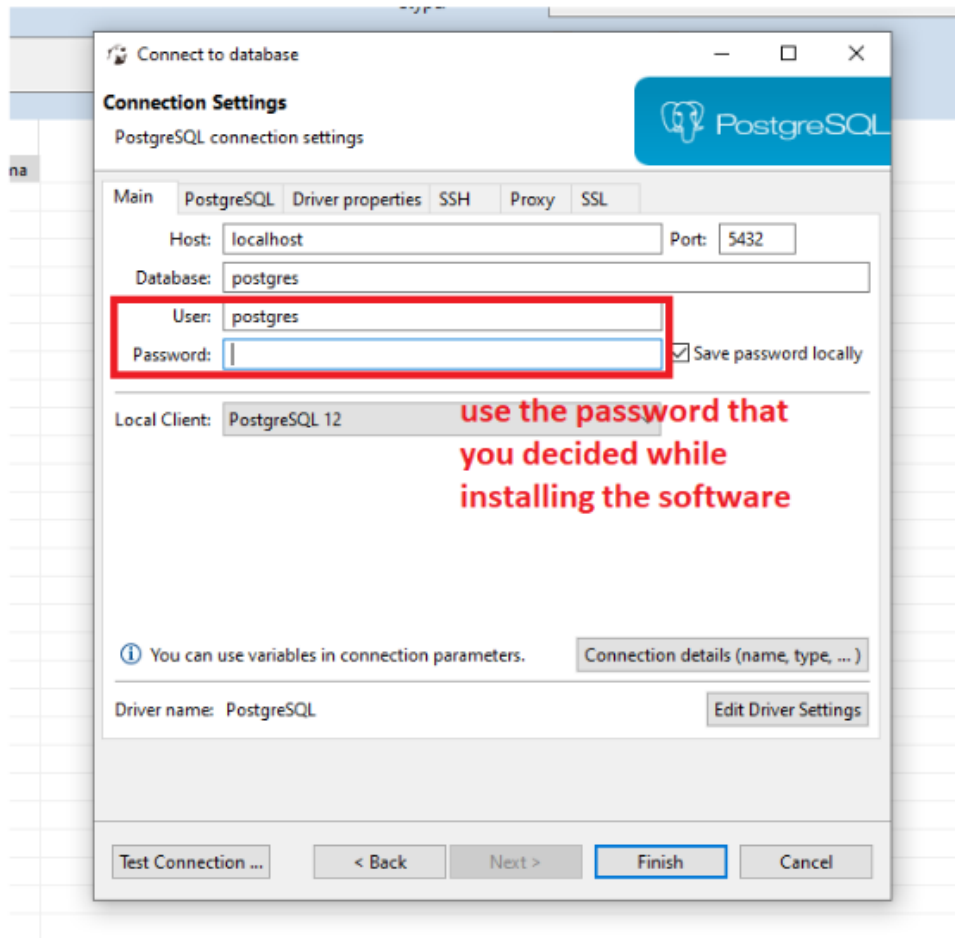


Figure 3: Enter credentials and connect

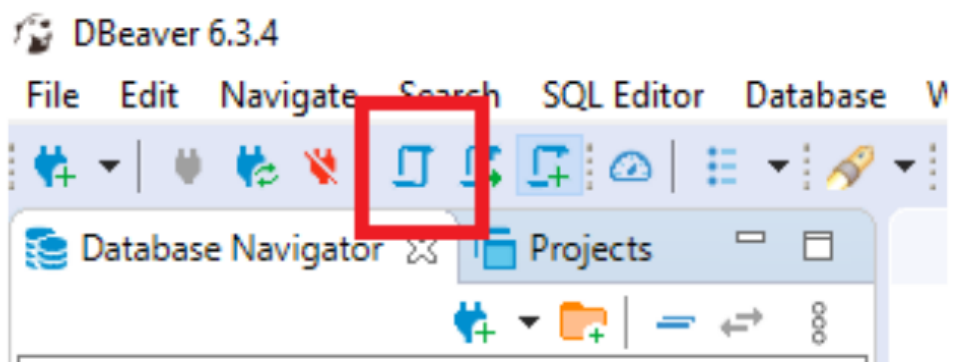


Figure 4: Open new script

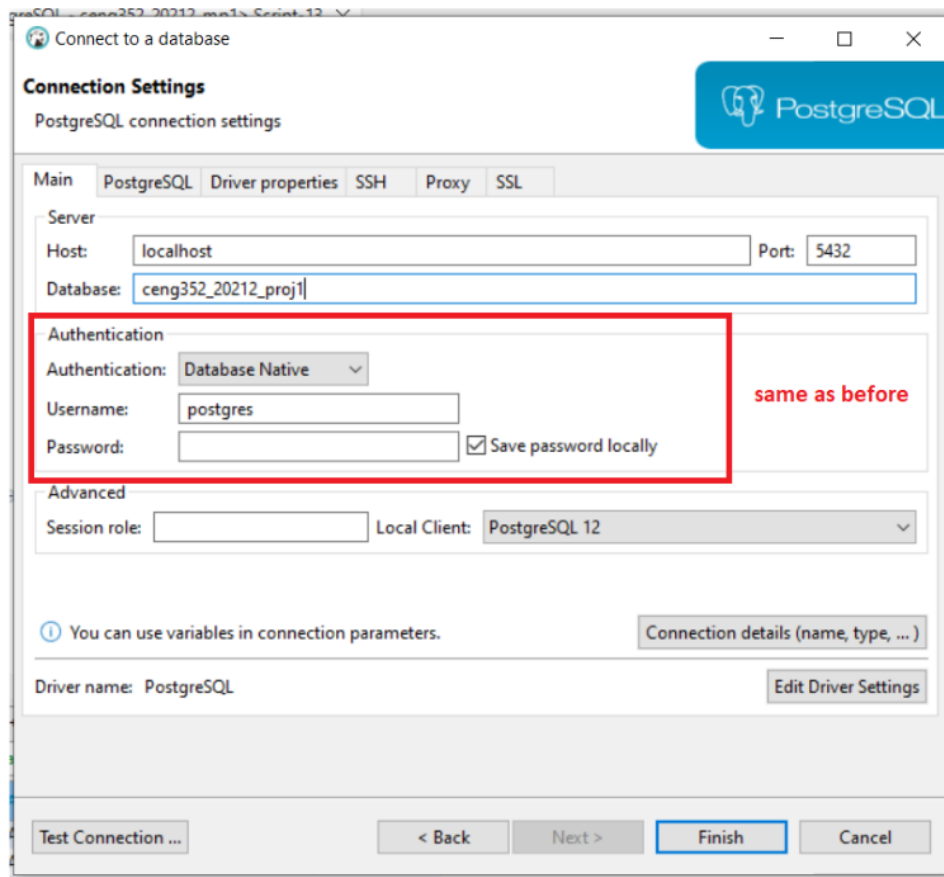


Figure 5: Connect to new database

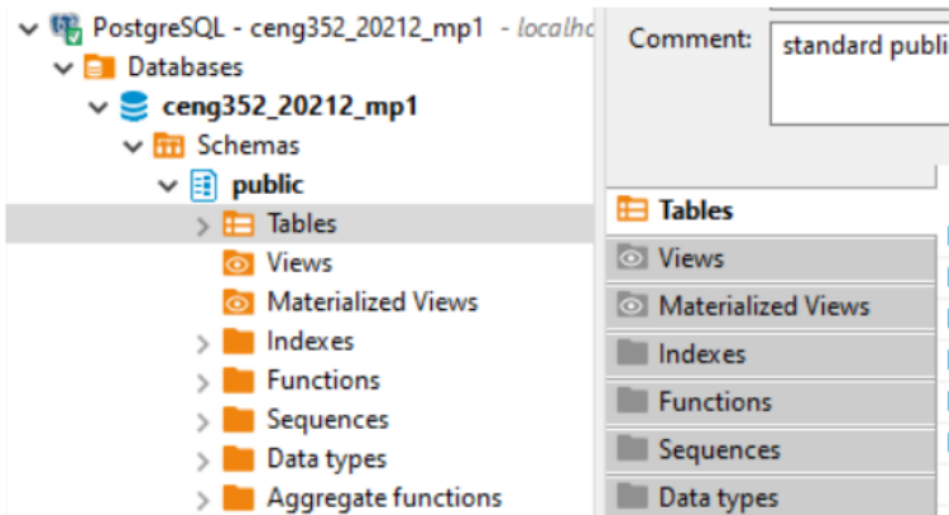


Figure 6: Open tables view

```

* <PostgreSQL - ceng352_20212_mp1> Script-13 public * <PostgreSQL - ceng352_20212_mp1> create_table_queries.sql X
create table if not exists actors (
    ...
    primary key ("actorId")
);

create table if not exists directors (
    ...
    primary key ("directorId")
);

create table if not exists movies (
    ...
    primary key ("movieId")
);

create table if not exists directed (
    ...
);

create table if not exists genres (
    ...
);

create table if not exists casts (
    ...
);

```

Figure 7: Run create scripts

Properties ER Diagram PostgreSQL - ceng352\_20212\_mp1

Name: public Namespace ID: 2200  
 Comment: standard public schema Owner: postgres

Tables	Table Name	Object ID	Owner	Tablespace	Row Count Estimate	Partitions	Partiti
	actors	127,363	postgres	pg_default	3,016,048	[ ]	
Views	cast	Create		default	1,585,794	[ ]	
Materialized Views	dire	View Table	F4	default	491,088	[ ]	
Indexes	dire	Filter		default	742,505	[ ]	
Functions	gen	View Diagram		default	625,353	[ ]	
Sequences	mov	View Data		default	511,480	[ ]	
Data types		Compare/Migrate					
Aggregate functions		Export Data					
Permissions		Import Data					
Source		Tools					
		Generate SQL					
		Read data in SQL editor					
		Copy	Ctrl+C				
		Paste	Ctrl+V				
		Copy Advanced Info	Ctrl+Shift+C				
		Delete	Delete				
		Rename	F2				
		Refresh	F5				

Figure 8: Import data

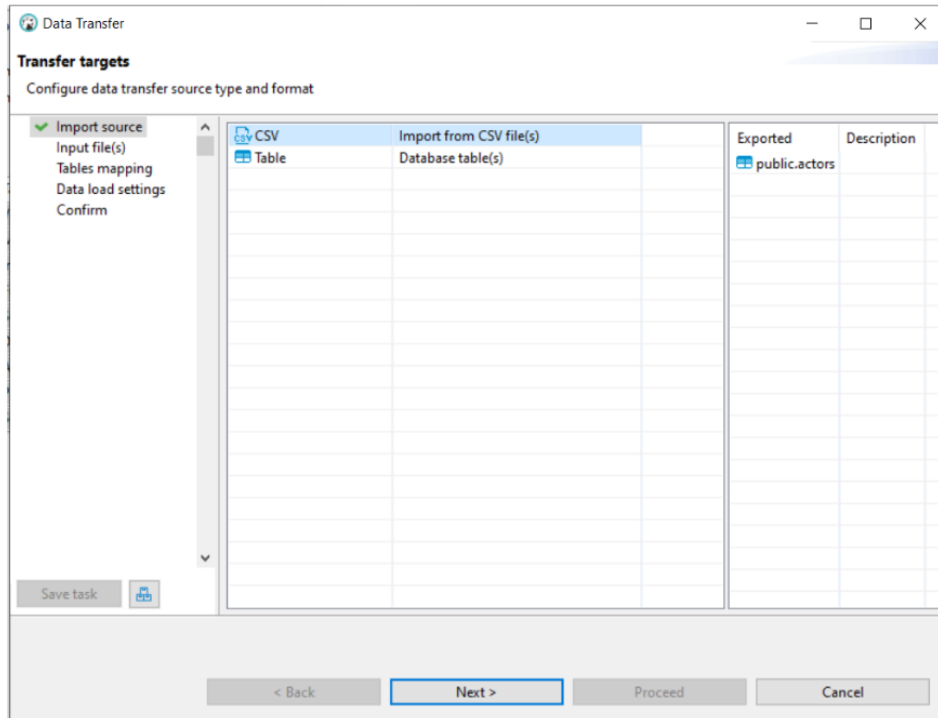


Figure 9: Choose CSV

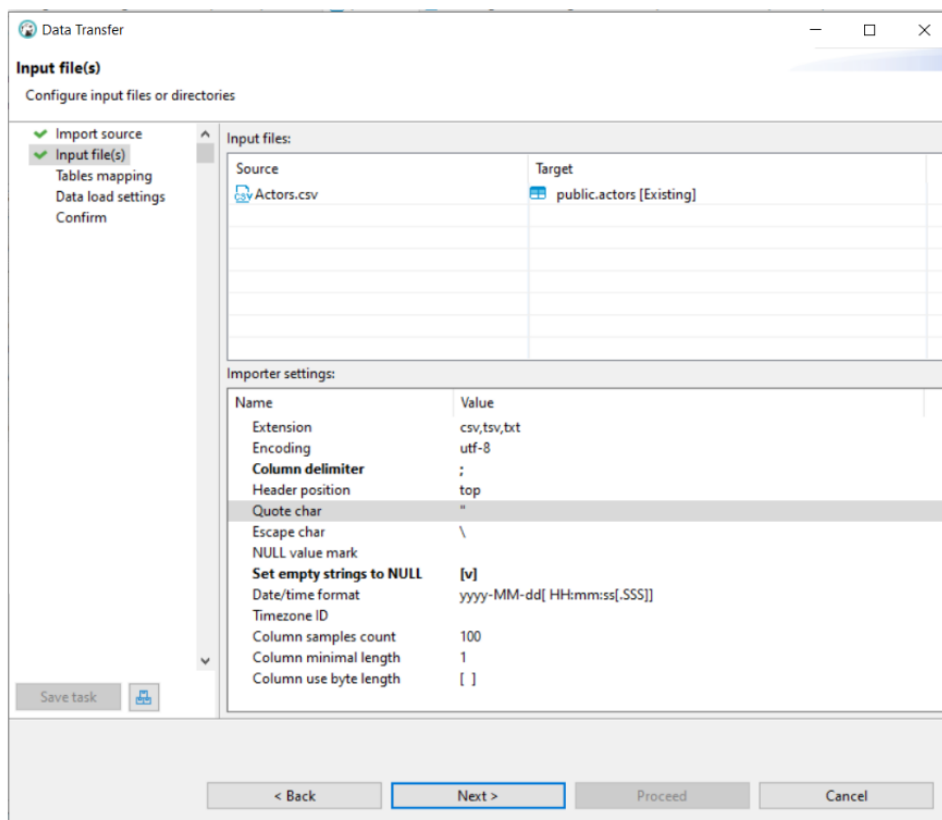


Figure 10: Set delimiter